



EMERITUS

Learn. From the world's best.

Paradigms of Machine Learning - II



Why Strategy Matters in ML



Paradigm Impact

Impacts data needs,
time, accuracy and trust



Resource Efficiency

Saves time and money
with the right choices



Success Rate

Boosts chances of
project success

Constraints That Affect ML Decisions

Factors That Shape Machine Learning Choices

Constraint	Examples
Data Volume	50 labelled records only
Label Cost	Expensive expert labeling
Privacy	Cannot send data to server
Compute	Runs on mobile device only
Domain Gap	Training vs. deployment mismatch


Learning Paradigms in AI

Different Ways AI Models Learn From Data




Supervised

Learns using labelled data



Self-Supervised

Generates its own labels from raw data




Transfer Learning

Uses knowledge from pre-trained models



Few-shot / One-shot

Learns from very few examples




Federated

Trains on multiple devices without sharing data


Important Questions to Consider

Factors That Guide the Choice of Learning Approach



Paradigm Trade-offs

Key Factors to Weigh When Choosing A Learning Approach



Learning Paradigms at a Glance

Quick Comparison of Learning Types, Data Needs and Use Cases

Paradigm	Labels Required	Training Data Size	Adaptability	Used In
Supervised	✔ High	✔ High	✘ Low	Spam filters, fraud detection
Unsupervised	✘ None	✔ High	✘ Low	Clustering, dimensionality reduction
Semi-supervised	⚠ Few	✔ High	⚠ Medium	Sentiment analysis with partial labels
Self-supervised	✘ None	✔ Very High	✔ High	Pre-training for NLP/vision models
Active Learning	✔ Selective	⚠ Medium	✔ Iterative	Medical imaging, rare labelling
Transfer Learning	✔ Pretrained	✘ Low (target)	✔ High	Language models, domain-specific tasks
Few-shot/One-shot	✔ 1–5 samples	✘ Very Low	✔ Extreme	Face ID, low-resource NLP
Federated Learning	✔ Distributed	✔ Localised	✔ Secure	Mobile keyboards, healthcare AI

Decision-Making Steps

Choose the Right Learning Approach Together

Share the Case

Outline needs and limits

Discuss Options

Compare suitable paradigms

Explain the Choice

Share logic and trade-offs



Case 1: Rural Health Assistant

Designing an AI tool For Remote Healthcare Support

Scenario

- Mobile health assistant for rural users
- Categorises symptoms as mild, moderate or severe to guide care

Challenge

- Use lightweight, low-resource models
- Ensure privacy with on-device data handling

Constraints


- Low data
- Strict privacy
- Basic devices

Activity: Choose the Right Approach

Which Method Fits the Challenge Best?



Transfer Learning: Use existing medical models



Federated Learning: Keep data on the device



Few-shot Learning: Train with very few examples

Case 2: Multilingual News Classifier

Detecting Bias and Misinformation in 12 Languages

Scenario

Build a system to classify news by topic and spot bias in 12 languages before elections

Constraints

- Very limited data
- Tight deadlines

Challenge

How to design a language-neutral model that works well with little data and scales quickly?

Activity: Which Approach Fits Best?

Choose The Right Method for Fow-Resource Language Tasks

- **Zero-shot**

Works without any training examples

- **Few-shot**

Learns from just a few examples

- **Translate + Supervised**

Translate to one language, then train normally

Case 3: Classifier for Remote Devices

Smart Image Detection on Low-Power, Offline Devices

Scenario

Use IoT devices (e.g. wildlife or farm cameras) to detect events like pests, crop stages or animals — all offline

Constraints

- Very limited processing power
- No internet access
- 2GB of local, unlabelled images

Learning Challenge

How do we learn from unlabelled images, stay lightweight and support smart local decisions?

Activity: What Would You Choose?

Pick an Approach and Explain Why



Self-supervised

Learn without labelled data



Transfer learning

Start with a pre-trained model



Justify

Explain reasoning



Discuss

Consider trade-offs

Demo Setup: Healthcare Feedback Classifier

Key Challenges

- **Problem:** Classify feedback as positive/negative
- **Few Labels:** Limited labelled examples available
- **Privacy Concern:** Must protect patient data
- **Low Compute:** Restricted processing power



Step 1: Load Pretrained Embeddings

Efficient Model Setup

HuggingFace Integration


Use pre-built models

Embedding Visualisation

Utilise existing knowledge

Implementation

Skip training from scratch



Step 2: Fine-tune on 100 Examples

Rapid Model Adaptation



Prepare Data: Small labelled dataset



Fine-Tune Model: Adapt pre-trained weights



Evaluate: Check performance



Iterate: Quick prototyping cycles

Step 3: Federated Learning Benefits

Privacy-Preserving AI Training

1

Local Training: Model trains on individual devices

Data Privacy: Patient data stays on device

2

3

Secure Sharing: Only model updates are shared

Central Aggregation: Server combines updates to improve model

4

Output Review and Discussion

Choosing Your Approach

Centralised Approach

- 87% accuracy | Faster training
- Privacy risks

Federated Approach

- 82% accuracy | Slower convergence
- Robust privacy protection

AI Toolkit Cheat Sheet

Essential Frameworks by Approach



Transfer Learning

HuggingFace | TensorFlow Hub



Few-shot Learning

PyTorch Lightning | OpenAI APIs



Federated Learning

Flower | TensorFlow Federated



Self-supervised

SimCLR | BYOL (PyTorch)

Paradigm-to-Constraint Visual Matrix: Summary

Constraint ↓ \ Paradigm →	Supervised	Transfer	Few-shot	Self-supervised	Federated
Low data availability	✗	✓	✓	⚠	✓
High label cost	✗	✓	✓	✓	✓
Strong privacy needs	✗	⚠	✓	✓	✓
Low compute budget	⚠	✓	✓	⚠	⚠

Common Mistakes When Choosing ML Paradigms

Pitfalls of Overusing Few-shot Learning

Be Mindful of When it Works — and When it Doesn't

Few-shot fails with imbalanced or noisy data

Leads to poor results when data isn't evenly distributed



Few-shot isn't a one-size-fits-all

Not ideal for all low-data scenarios

Misusing Transfer Learning

When Pretraining Goes Wrong



Wrong Source: Pretrained model from an unrelated domain



Poor Transfer: Knowledge doesn't apply well



Bad Outcome: Leads to negative transfer

Assuming all pretrained models improve performance without validation

Misunderstanding Federated Learning


Overlooking Key Risks and Limitations

Security Assumptions:

Not all federated setups are safe from attacks

Performance Limits:

Missed impact of communication and sync delays



Misusing Self-Supervised Learning

Key Pitfalls to Avoid

Undefined Pretext Task: Applying without a clear pretext task



Meaningless Embeddings: Leads to representations that fail to capture useful patterns



Efficiency Assumptions: Assuming it's always more efficient than supervised learning



General Pitfall: Ignoring Constraints


Key Oversights to Avoid

- **Data Availability:** Failing to consider data quantity and quality
- **Buzzword-Driven Choices:** Choosing methods based on trends, not suitability
- **Compute Resources:** Overlooking hardware and processing limits
- **Privacy/Legal Boundaries:** Ignoring regulatory and ethical restrictions



Impact of Paradigm Pitfalls

Consequences of Misapplication



Project Failure: Inability to achieve objectives



Resource Waste: Excessive time and budget spent



Suboptimal Performance: Models underperforming against their potential

Misapplying machine learning paradigms leads to these negative outcomes, especially when based on buzzwords rather than data availability, resources and legal/privacy constraints

Best Practices for Paradigm Selection

Making Informed Choices

Thorough Analysis

Evaluate your use case, data, and constraints carefully

Validation Testing

Test assumptions with small experiments before full rollout

Hybrid Approaches

Combine paradigms to balance strengths and weaknesses

Avoiding common pitfalls means focusing on practical needs, not trends. Understand the limits of each paradigm—few-shot, transfer, federated and self-supervised—to choose wisely

Real-World Industry Snapshots

Federated Learning in Action

How Google Gboard Learns Without Compromising Privacy



On-device word prediction

Trains locally on user typing data



Improves suggestions without sharing keystrokes

Preserves user privacy while enhancing performance




Why it Fits

Decentralised data and strict privacy requirements make federated learning ideal

Few-shot / Zero-shot Learning with GPT

Do more with fewer examples



Used for chatbots, tagging, coding help

Works across many tasks



Needs just 1–3 examples

For text generation or classification



Why it works

Fast to deploy, low data needed

Transfer Learning in Medical Imaging

Faster AI Solutions for Healthcare Startups

Use pretrained models

ResNet or EfficientNet on ImageNet

Fine-tune with 500–1000 scans

X-rays or MRIs for medical use

Why it works

Cuts labelling cost and speeds up launch



Self-Supervised Learning with DINO

Powerful Vision Models Without Labelled Data

Vision Transformer without labels:

Advanced computer vision without manual annotation



Learns from unlabelled images:

Extracts meaningful patterns autonomously



For segmentation and anomalies:

Practical applications in various domains



Why it fits:

Unlabelled data is plenty, labels are costly



Domain Adaptation in E-commerce

Adapting Review Models Across Regions

Cross-Region Adaptation

- U.S. models adapted for UK, India, etc.
- Core function kept, regional tweaks applied

Contextual Adjustments

- Adapts tone, product terms and behaviour
- Detects cultural and language nuances

Strategic Advantage

- Reuses core model with minor tweaks
- Efficient for global e-commerce

Q&A Session



Open floor for questions and clarifications.